



Trigger functions

Zabbix supports almost 100 different functions:

⚡ Aggregate functions	avg, count, min, max, sum, kurtosis etc.
⚡ Bitwise functions	bitand, bitlshift, bitnot, bitor, bitrshift, bitxor
⚡ Date and time functions	date, dayofmonth, dayofweek, now, time
⚡ History functions	last, first, change, logseverity, monoinc, nodata, etc.
⚡ Trend-based functions	trendavg, trendcount, trendmax, trendmin, trendsum
⚡ Mathematical functions	abs, cos, sin, tan, ceil, floor, degrees, e, exp, log, rand, etc.
⚡ Operator functions	in, between
⚡ Prediction functions	forecast, timeleft
⚡ String functions	ascii, bitlength, char, concat, find, left, length, trim, mid etc.
⚡ Foreach functions	avg_foreach, last_foreach, sum_foreach (calculated items only)

Advanced functions and absolute time shift periods are discussed in ZCP 6.0

Most of the functions require one or multiple parameters:

⚡ `/host/key` is a common mandatory first parameter for history functions

```
last(/prod/system.cpu.load)
```

⚡ Other parameters are placed after the `/host/key` separated by a comma

⚡ If an evaluation period or range is required, it always goes as a second parameter

```
min(/prod/vm.memory.size[free],1h)
```

⚡ More than one parameter may be required for some functions

```
count(/prod/log[/var/log/myApp.log],10m,"like","error" )
```



`/host/key` and evaluation period or range parameters must never be quoted


Most of numeric functions accept time or number of values as a parameter:

- ⚡ Seconds will be used if **no time suffix** is specified
- ⚡ Time suffixes may be used to specify **time units** (10s, 5m, 1h, etc.)
- ⚡ If preceded by a hashtag #, the parameter indicates the number of values (#5, #10, etc.)

sum(/host/key,10m)	➡	Sum of values in the last 10 minutes
sum(/host/key,10s)	➡	Sum of values in the last 10 seconds
sum(/host/key,10)	➡	Sum of values in the last 10 seconds
sum(/host/key,#10)	➡	Sum of the last 10 values

Example:

- Item update interval: 5s
- Item collected every minute at 00, 05, 10, 15 ,etc.



Mathematical operations can be applied to the trigger functions:

⚡ To the result of trigger function

`avg(/host/net.if.in[eth0,bytes],10m) * 8 > 10M`

result is multiplied by 8

⚡ Between results of trigger functions

`min(/host/system.cpu.load,5m) / last(/host/system.cpu.num) > 1.5`

Load per CPU

⚡ To the results of multiple trigger functions

`last(/host/proc.num) / last(/host/kernel.maxproc) * 100 > 80`

80 percent calculated

Mathematical operations must be added after the function

Correct

`min(/host/key,#5)`

Correct

`min(/host/key,#5)*10`

Wrong

`min(/host/key,#5*10)`

Aggregate functions will accept other expressions as function parameters:

`function(function_a(),function_b(),...)`

⚡ **Function** beginning with **host** and **item key** will use **time period** as parameter

`min(/host/key1,1h)`

smallest value from 1 hour of historical data

⚡ **Function** beginning with **other expressions** will use them as parameters

`min(avg(/host/key1,1h),min(/host/key2,#5)*10)`

smallest value from the result of other expressions



function1



function2

The following operators are supported for triggers:

- ⚡ Unary minus - (change the sign of an operand)
- ⚡ Mathematical operations (+, -, *, /)
- ⚡ Compare (<, <=, >, >=, =, <>)
- ⚡ Logical operators (and, or, not)
 - ✓ Case-sensitive and must be in lowercase
 - ✓ Must be surrounded by spaces or parentheses.

Notes:

- ⚡ Most operators expect numerical variables
- ⚡ Operators = or <> can be used to compare strings

`last(/host/vfs.file.cksum[/etc/passwd],#1) <> last(/host/vfs.file.cksum[/etc/passwd],#2)`

`last(/host1/system.hw.macaddr[eth0,short],#1) = last(/host2/system.hw.macaddr[eth0,short],#1)`

When building trigger expressions, it is possible to compare function against:

⚡ Fixed value

`last(/prod/agent.version) <> '6.0.0'`

⚡ User macro

`min(/prod/system.cpu.load,5m) > {$CPU.LOAD}`

⚡ Result of another trigger function

`last(/node1/hw.macaddr) = last(/node2/hw.macaddr)`

⚡ Result of calculation

`last(/node1/system.cpu.num) > last(/node2/system.cpu.num) * 1.5`

An optional time shift is supported in the time parameter:

⚡ This parameter allows to reference the data from a period of time in the past


⚡ Time shift starts with now (the current time), followed by:

- ✓ +N<time unit> to add N time units
- ✓ -N<time unit> to subtract N time units

`avg(/host/key, 4h:now-1d)`

⚡ Complex expressions using multiple time units in the calculation are supported

`avg(/host/key, 6h:now-1d+8h)`



<https://www.zabbix.com/documentation/6.0/en/manual/config/triggers/expression>

The last() function returns the **last received value**:

⚡ The time period is not supported

`last(/host/key)`

⚡ Using hashtag # will denote the Nth previous value

`last(/host/key,#1)`



Last value

`last(/host/key,#3)`



3rd previous value

⚡ Time shift parameters are supported

`last(/host/key,#1:now-1h)`

Value received one hour ago

`last(/host/key,#4)`

`last(/host/key,#3)`

`last(/host/key,#2)`


`last(/host/key)`
`last(/host/key,#1)`

time




Function `last()` is very sensitive:

- ⚡ Every received value which exceeds the threshold will generate a new alert
- ⚡ This can lead to "trigger flapping"



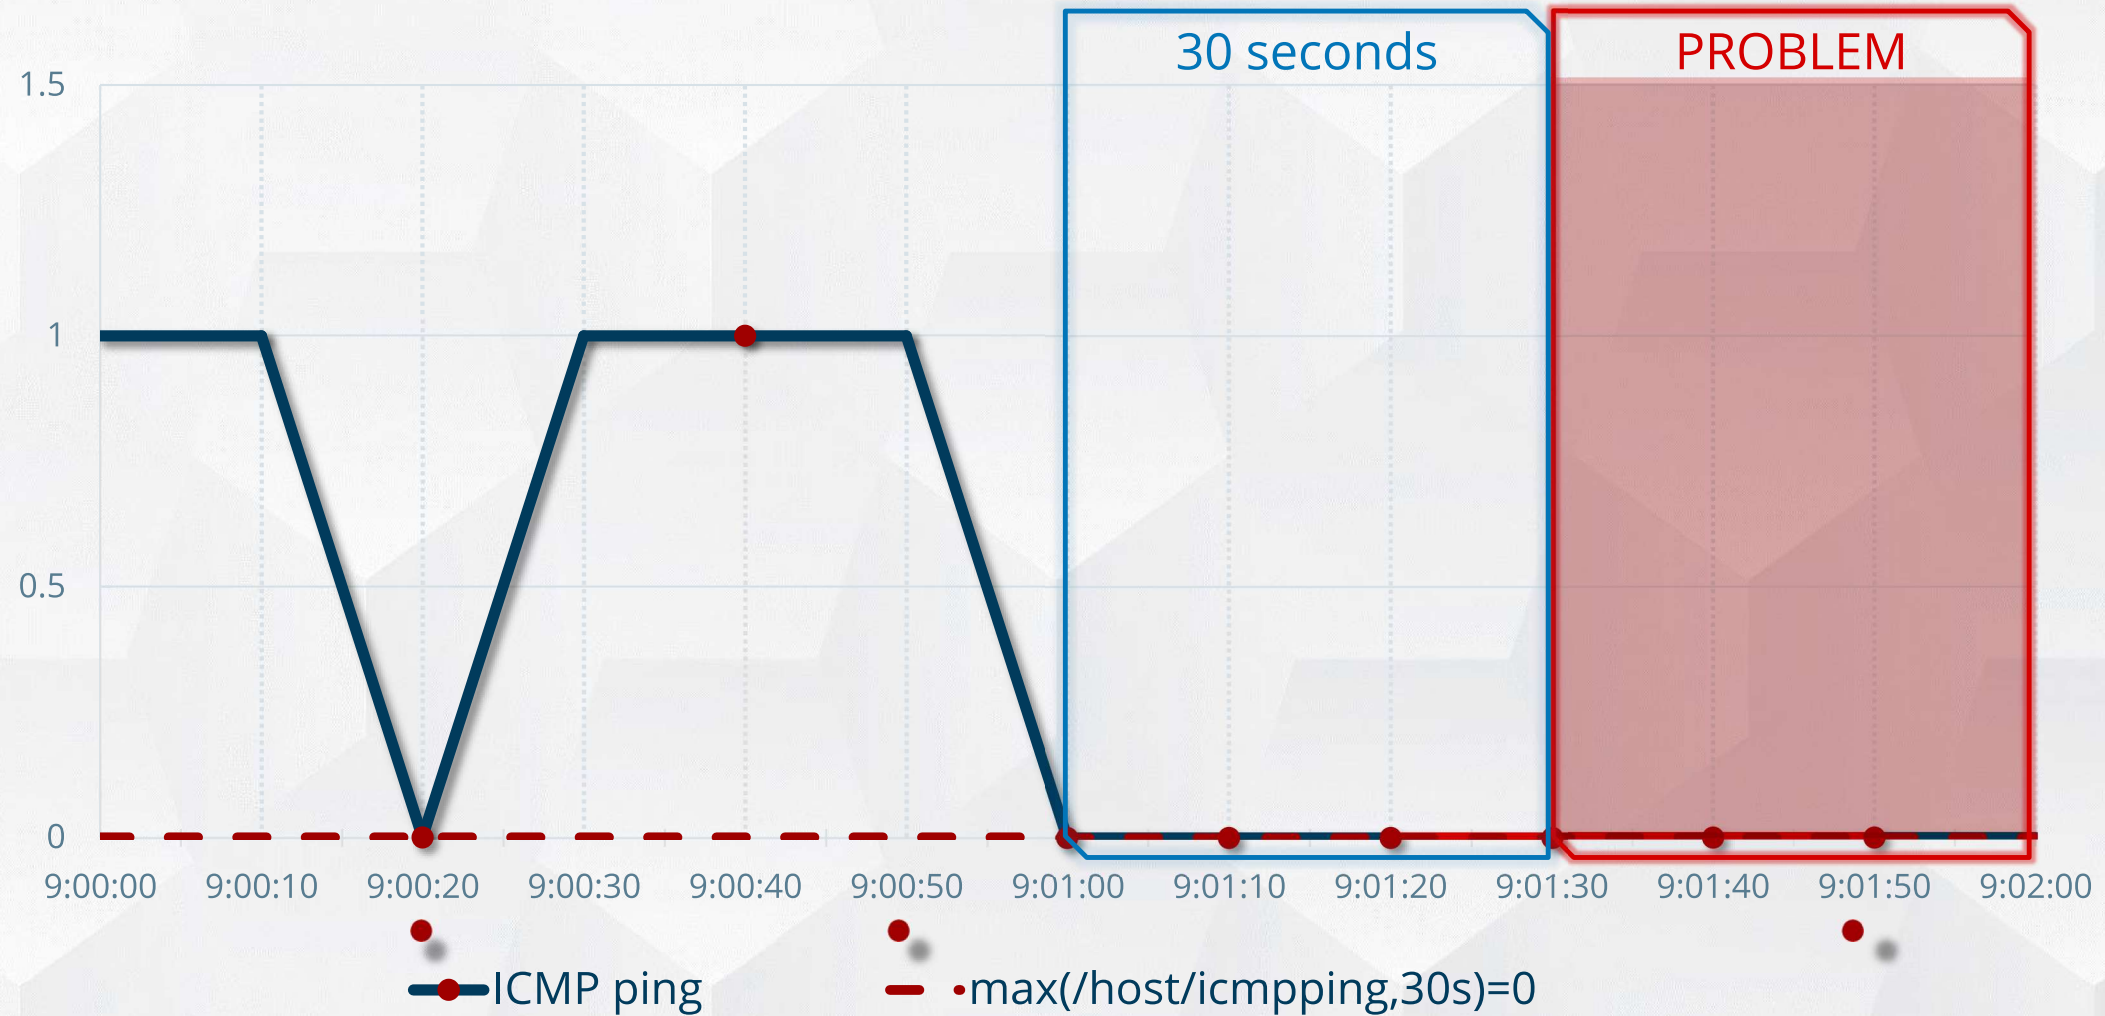
Function min() is a simple way to reduce false problem detection:

- ⚡ Minimum value must drop below the threshold for some time period to detect a problem
- ⚡ The problem will resolve immediately




Function max() may be used to detect availability issues:

⚡ Multiple availability checks in a row must fail to detect a problem




Function count() is another option:

⚡ More than one anomaly will be required to detect a problem



Monotonic functions can be used to monitor queues or disk space:

- ⚡ monoinc() detects **monotonic increase** of values collected
- ⚡ monodec() detects **monotonic decrease** of values collected



PRACTICAL SETUP


- 1) On the host "Training-VM-XX":
 - ✓ Replace last() function for CPU load triggers with a minimum for 1 minute
- 2) Use "cat /dev/urandom | md5sum" command to test triggers
- 3) Create a new trigger to compare memory usage:
 - ✓ Compare average free memory for current hour with average free memory for previous hour
 - ✓ Generate alert if the free memory has decreased by more than 25%



Time-based functions

All **time-based functions** are recalculated **every 30 seconds**:

- ⚡ If both time-based and non-time-based functions are used in an expression, it is recalculated when a new value is received and additionally every 30 seconds
- ⚡ Recalculation schedule is distributed evenly between all time-based functions
- ⚡ Example trigger is recalculated:
 - ✓ every 10 seconds based on the item update interval
 - ✓ additionally, every 30 seconds because time-based function is used in the expression



All **date and time functions** are time-based:

📶 date	current date in YYYYMMDD format
📶 time	current time in HHMMSS format
📶 dayofweek	current day of week 1-7 (Mon - 1, Sun - 7)
📶 dayofmonth	current day of month 1-31
📶 now	current time in HHMMSS format

These functions can be used to specify time periods from trigger calculation

📶 Detect problems only on weekends

`min(/host/system.cpu.load,30s) > 2 and dayofweek() > 5`

📶 Ignore scheduled backups between 01:00:00 and 03:00:00

`avg(/host/system.cpu.util[,iowait],5m) > 5 and (time() < 010000 or time() > 030000)`




All date and time functions use Zabbix server time zone!

Time-based history function checks for no data received:

`nodata(/host/key,time period,<mode>)`


- ⚡ Time period should **not be less than 30 seconds** (`nodata(0)` is not allowed)
- ⚡ Returns:
 - ✓ 1 - if no data received during the defined period of time
 - ✓ 0 - otherwise
- ⚡ The 'nodata' triggers monitored by proxy are, by default, **sensitive to proxy availability**
 - ✓ They will not fire if the data is expected from a proxy, which is currently offline
 - ✓ "strict" mode will ignore proxy availability



Function `nodata()` can be used to detect:

- ⚡ Zabbix agent availability (agent.ping item never returns 0)
- ⚡ Changes in the log files monitored by Zabbix
- ⚡ Data received (or not received) on the regular intervals

The problem is resolved immediately when the data arrives



Zabbix server time zone is used to calculate time-based functions

- ⚡ User time zone settings may differ from Zabbix server time zone

Time-based functions in triggers with **multiple event generation** mode will create a new problem every 30 seconds

- ⚡ The trigger will be evaluated every 30 seconds even if there isn't any new data received


Nodata() function may return **false positives** if:

- ⚡ There are time differences between Zabbix server, proxy and agent
- ⚡ Discard unchanged preprocessing steps are used
- ⚡ History is not saved for the item

Nodata() function is evaluated for "not supported" items also

PRACTICAL SETUP

- 1) On the "Training-VM-XX active checks" host
 - ✓ Create Zabbix agent ping item with the 10 second update interval
 - ✓ Create a trigger to check agent ping last value
- 2) Stop Zabbix agent on your virtual machine
- 3) Wait 1 minute to test if the trigger detects a problem
- 4) On the Training-VM-XX active checks host:
 - ✓ Replace the last() trigger function with no data received for 1 minute
- 5) Wait 1 minute to test if the trigger detects a problem
- 6) Start Zabbix agent on your virtual machine




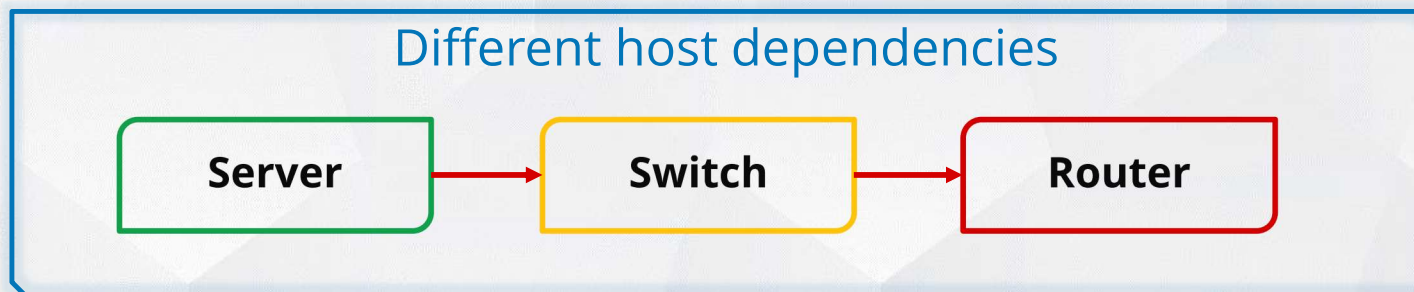
TRIGGER DEPENDENCIES

Dependencies can be defined between multiple triggers:

- ⚡ Problems will be suppressed if the trigger they depend on is in the PROBLEM state
- ⚡ Zabbix does not support dependencies between hosts directly

Dependencies between triggers can be defined:

- ⚡ On the same host:
 - ✓ Problem level (different severities)
- ⚡ Different hosts:
 - ✓ Network devices
 - ✓ Applications
 - ✓ Other resources



It is possible to create complex architecture with multi-level dependency:


⚡ Multiple levels:

✓ Server A > Switch 1 > Router

⚡ Multiple dependencies:

✓ Switch 1 > Router

✓ Switch 2 > Router



If, for example, a router is down, and dependencies are defined:

⚡ Problems generated by the dependent triggers will be **suppressed** and hidden

⚡ Zabbix will not execute actions for the dependent trigger

⚡ The dependent trigger will be re-evaluated and will change its state only after the parent trigger returns to the OK state and new metrics are received.

Dependent trigger will only be re-evaluated when:

- ⚡ Parent trigger has changed its state to "OK"
- ⚡ A new value is received for item used in dependent trigger expression

Triggers are evaluated independently of their dependencies:

- ⚡ It is possible that problem with a dependent trigger will be detected first
- ⚡ In this scenario, the dependent trigger will fire as usual
 - ✓ It will become suppressed later when problem with a parent trigger will be detected

PRACTICAL SETUP

- 1) On the host "Training-VM-XX":
 - ✓ Create user macro for very high CPU load (>2.5) with a "High" severity
 - ✓ Create a trigger to detect very high CPU load
- 2) Create a dependency between triggers for high and very high CPU load
- 3) Use "yes > /dev/null" command to test this setup